

```
import nltk
from sklearn.feature_extraction.text import TfidfVectorizer
nltk.download('stopwords')
from tensorflow.keras.preprocessing.text import Tokenizer

from sklearn.preprocessing import LabelEncoder

[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data]   Package stopwords is already up-to-date!

import tensorflow as tf
from sklearn.model_selection import train_test_split
from tensorflow.keras import datasets, layers, models, preprocessing
from keras.models import Sequential
from keras.layers import Dense, SimpleRNN
from nltk.corpus import stopwords
import pandas as pd
import numpy as np
import seaborn as sb

df = pd.read_csv('youtoxic_english_1000.csv')
labels = []
y = df.IsToxic
df_y = pd.DataFrame(y, columns=['IsToxic'])

sb.catplot(x="IsToxic", kind='count', data=df_y)


i = np.random.rand(len(df)) < 0.8
train = df[i]
test = df[~i]
num_labels = 2
vocab_size = 25000
batch_size = 100

tokenizer = Tokenizer(num_words=vocab_size)
tokenizer.fit_on_texts(train.Text)


x_train = tokenizer.texts_to_matrix(train.Text, mode='tfidf')
x_test = tokenizer.texts_to_matrix(test.Text, mode='tfidf')
df.apply(LabelEncoder().fit_transform)

encoder = LabelEncoder()
```

! 0s completed at 8:39 PM



```
y_train = encoder.transform(train.IsToxic)
y_test = encoder.transform(test.IsToxic)
```



<https://www.kaggle.com/datasets/reihanenamdari/youtube-toxicity-data?resource=download>

This dataset contains 1000 comments with labels of IsToxic, IsAbusive, IsThreat, IsProvocative, IsObscene, IsHatespeech, IsRacist, IsNationalist, sSexist, IsHomophobic, IsReligiousHate, IsRadicalism. Despite the small dataset, most of the Text in the dataset is fairly long.

In the collab script. We first start by loading in the dataset using pandas. I went ahead and plotted for all 11 columns, but there was a lot of invalid fields within those columns. So I had to do some data cleaning for those fields. After plotting, we could see that IsToxic has the most count out of all of the comments, followed by IsAbusive.

```
-----
ValueError                                Traceback (most recent call last)
<ipython-input-37-258ae25c3c08> in <module>
----> 1 sb.catplot(x="IsThreat", kind='count', data=pd.DataFrame(y, columns=
['IsThreat']))
```

2 frames

```
/usr/local/lib/python3.8/dist-packages/seaborn/categorical.py in
establish_colors(self, color, palette, saturation)
217         # Determine the gray color to use for the lines framing the plot
```

```

317         # determine the gray color to use for the lines framing the plot
318         light_vals = [colorsys.rgb_to_hls(*c)[1] for c in rgb_colors]
--> 319         lum = min(light_vals) * .6
320         gray = mpl.colors.rgb2hex((lum, lum, lum))
321

```

**ValueError:** min() arg is an empty sequence

```

model = models.Sequential()
model.add(layers.Dense(16, input_dim=vocab_size, kernel_initializer='normal', activation='
#model.add(layers.Dense(8, activation='relu'))

model.add(layers.Dense(1, kernel_initializer='normal', activation='sigmoid'))
model.compile(loss='binary_crossentropy',
              optimizer='adam',
              metrics=['accuracy'])

```

```

from keras.utils import np_utils

```

```

#y = np_utils.to_categorical(y_train, 1)
#y = np_utils.to_categorical(y_train, 2)
print(x_train.shape)
#WOY = np_utils.to_categorical(OY,5)

```

```

#y_trainEnc = y.sort_indices()
#y_testEnc = OY.sort_indices()

```

```

model.summary()

```

```

Seq = model.fit(x_train,y_train,batch_size=batch_size,epochs=15,validation_split=0.1,verbo
#model.fit(x,y,batch_size=32,epochs=1,callbacks=callbacks,validation_data=(OX,OY))

```

```

(796, 25000)
Model: "sequential_6"

```

Layer (type)	Output Shape	Param #
=====		
dense_13 (Dense)	(None, 16)	400016
dense_14 (Dense)	(None, 1)	17

```

=====
Total params: 400,033
Trainable params: 400,033
Non-trainable params: 0

```

```

Epoch 1/15

```

```
Epoch 1/15
8/8 [=====] - 2s 43ms/step - loss: 0.6944 - accuracy: 0.4874
Epoch 2/15
8/8 [=====] - 0s 21ms/step - loss: 0.6687 - accuracy: 0.6591
Epoch 3/15
8/8 [=====] - 0s 18ms/step - loss: 0.6436 - accuracy: 0.7500
Epoch 4/15
8/8 [=====] - 0s 20ms/step - loss: 0.6136 - accuracy: 0.8191
Epoch 5/15
8/8 [=====] - 0s 22ms/step - loss: 0.5763 - accuracy: 0.8711
Epoch 6/15
8/8 [=====] - 0s 20ms/step - loss: 0.5339 - accuracy: 0.9061
Epoch 7/15
8/8 [=====] - 0s 21ms/step - loss: 0.4875 - accuracy: 0.9281
Epoch 8/15
8/8 [=====] - 0s 22ms/step - loss: 0.4410 - accuracy: 0.9461
Epoch 9/15
8/8 [=====] - 0s 20ms/step - loss: 0.3961 - accuracy: 0.9601
Epoch 10/15
8/8 [=====] - 0s 21ms/step - loss: 0.3543 - accuracy: 0.9691
Epoch 11/15
8/8 [=====] - 0s 22ms/step - loss: 0.3166 - accuracy: 0.9741
Epoch 12/15
8/8 [=====] - 0s 22ms/step - loss: 0.2828 - accuracy: 0.9741
Epoch 13/15
8/8 [=====] - 0s 25ms/step - loss: 0.2532 - accuracy: 0.9801
Epoch 14/15
8/8 [=====] - 0s 18ms/step - loss: 0.2269 - accuracy: 0.9811
Epoch 15/15
8/8 [=====] - 0s 21ms/step - loss: 0.2040 - accuracy: 0.9831
```

```
score = model.evaluate(x_test, y_test, batch_size=batch_size, verbose=1)
```

```
print('Accuracy: ', score[1])
```

```
3/3 [=====] - 0s 6ms/step - loss: 0.6371 - accuracy: 0.6814
Accuracy: 0.6813725233078003
```

I went ahead and did the Sequential model, but Multi-Label classification is a bit more difficult than I intended, so I scraped all of the columns and stuck with `IsToxic` as the only target function. The Sequential performed extremely well, at first I had 4 layers ranging from 64, 32, 16, 1, which training accuracy was really well, however the evaluation was very low, maybe due to overfitting. I also tried some loss function, activation functions, and optimizer, which the `binary_crossentropy`, `sigmoid` and `rmsprop`, performed the best. After a few more run, I settled on a 16, 1 layer group which had the best accuracy of .68, a little disappointing but no matter.

```
model1 = Sequential()
model1.add(layers.Embedding(1000, 32))
model1.add(layers.SimpleRNN(32))
```

```

model1.add(layers.SimpleRNN(32))
model1.add(layers.Dense(1, activation='sigmoid'))
model1.compile(loss='binary_crossentropy',
               optimizer='rmsprop',
               metrics=['accuracy'])

```

```

batch_size = 32
maxlen = 800
train_data = preprocessing.sequence.pad_sequences(x_train, maxlen=maxlen)
test_data = preprocessing.sequence.pad_sequences(x_test, maxlen=maxlen)
model1.summary()

```

```
RNN = model1.fit(train_data,y_train,batch_size=batch_size,epochs=10,validation_split=0.2,v
```

Model: "sequential\_4"

Layer (type)	Output Shape	Param #
=====		
embedding_3 (Embedding)	(None, None, 32)	32000
simple_rnn (SimpleRNN)	(None, 32)	2080
dense_10 (Dense)	(None, 1)	33

```

=====
Total params: 34,113
Trainable params: 34,113
Non-trainable params: 0

```

```

Epoch 1/10
20/20 [=====] - 6s 213ms/step - loss: 0.6944 - accuracy: 0.4
Epoch 2/10
20/20 [=====] - 4s 201ms/step - loss: 0.6930 - accuracy: 0.4
Epoch 3/10
20/20 [=====] - 4s 199ms/step - loss: 0.6937 - accuracy: 0.4
Epoch 4/10
20/20 [=====] - 4s 199ms/step - loss: 0.6941 - accuracy: 0.4
Epoch 5/10
20/20 [=====] - 4s 197ms/step - loss: 0.6941 - accuracy: 0.4
Epoch 6/10
20/20 [=====] - 4s 199ms/step - loss: 0.6975 - accuracy: 0.4
Epoch 7/10
20/20 [=====] - 4s 197ms/step - loss: 0.6938 - accuracy: 0.4
Epoch 8/10
20/20 [=====] - 4s 197ms/step - loss: 0.6950 - accuracy: 0.4
Epoch 9/10
20/20 [=====] - 5s 273ms/step - loss: 0.6934 - accuracy: 0.4
Epoch 10/10
20/20 [=====] - 4s 198ms/step - loss: 0.6957 - accuracy: 0.4

```

```
from sklearn.metrics import classification_report
```

```

pred = model1.predict(test_data)
pred = [1.0 if p>= 0.5 else 0.0 for p in pred]
print(classification_report(y_test, pred))

```

```

7/7 [=====] - 1s 46ms/step
              precision    recall  f1-score   support

     0       0.55         1.00         0.71         115
     1       0.00         0.00         0.00          94

 accuracy                   0.55         209
 macro avg       0.28         0.50         0.35         209
 weighted avg    0.30         0.55         0.39         209

```

```

/usr/local/lib/python3.8/dist-packages/sklearn/metrics/_classification.py:1318: Under-
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.8/dist-packages/sklearn/metrics/_classification.py:1318: Under-
_warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.8/dist-packages/sklearn/metrics/_classification.py:1318: Under-
_warn_prf(average, modifier, msg_start, len(result))

```

After that is the RNN model, I went with an Embedding layer, then the RNN layer of size 32 then to the output layer. I decided the max len for each step of the RNN would be 15 as I believe after 15, there is a good part you could understand if the sentence is toxic enough, but also I didn't want to slow down the RNN model. In the end RNN models perform very poorly with an accuracy rate of .55. Most likely due to the size of the dataset.

```

train_data = preprocessing.sequence.pad_sequences(x_train, maxlen=15)
test_data = preprocessing.sequence.pad_sequences(x_test, maxlen=15)

```

```

model2 = models.Sequential()
model2.add(layers.Embedding(5000, 8, input_length=15))
model2.add(layers.Flatten())
model2.add(layers.Dense(32, activation='relu'))
model2.add(layers.Dense(8, activation='relu'))
model2.add(layers.Dense(1, activation='sigmoid'))
model2.compile(optimizer='rmsprop', loss='binary_crossentropy', metrics=['acc'])
model2.summary()

```

```

Emb = model2.fit(train_data, y_train, epochs=10, batch_size=32, validation_split=0.2)

```

```

Model: "sequential_3"

```

Layer (type)	Output Shape	Param #
--------------	--------------	---------

```

=====
embedding_2 (Embedding)      (None, 15, 8)              40000
flatten_2 (Flatten)          (None, 120)                 0
dense_7 (Dense)               (None, 32)                 3872
dense_8 (Dense)               (None, 8)                  264
dense_9 (Dense)               (None, 1)                   9

=====
Total params: 44,145
Trainable params: 44,145
Non-trainable params: 0

```

```

Epoch 1/10
20/20 [=====] - 1s 13ms/step - loss: 0.6938 - acc: 0.5016 - 
Epoch 2/10
20/20 [=====] - 0s 5ms/step - loss: 0.6936 - acc: 0.4921 - 
Epoch 3/10
20/20 [=====] - 0s 5ms/step - loss: 0.6933 - acc: 0.5047 - 
Epoch 4/10
20/20 [=====] - 0s 5ms/step - loss: 0.6932 - acc: 0.5111 - 
Epoch 5/10
20/20 [=====] - 0s 5ms/step - loss: 0.6932 - acc: 0.5111 - 
Epoch 6/10
20/20 [=====] - 0s 5ms/step - loss: 0.6931 - acc: 0.5111 - 
Epoch 7/10
20/20 [=====] - 0s 5ms/step - loss: 0.6931 - acc: 0.5111 - 
Epoch 8/10
20/20 [=====] - 0s 4ms/step - loss: 0.6933 - acc: 0.5111 - 
Epoch 9/10
20/20 [=====] - 0s 4ms/step - loss: 0.6931 - acc: 0.5111 - 
Epoch 10/10
20/20 [=====] - 0s 5ms/step - loss: 0.6932 - acc: 0.5111 - 

```

```

#score = model2.evaluate(test_data, y_test, batch_size=32, verbose=1)
pred = model2.predict(test_data)
pred = [1.0 if p>= 0.5 else 0.0 for p in pred]
print(classification_report(y_test, pred))

```

```

-----
ValueError                                Traceback (most recent call last)
<ipython-input-58-393c43ec3065> in <module>
      1 #score = model2.evaluate(test_data, y_test, batch_size=32, verbose=1)
----> 2 pred = model2.predict(x_train)

```

1 frames

```

/usr/local/lib/python3.8/dist-packages/keras/engine/training.py in
tf__predict_function(iterator)
     13         try:
     14             do_return = True
     15             actual = self._converted_call_fn(self._predict_function,

```

```

---> 15             retval_ = ag__.converted_call(ag__.ld(step_function),
(ag__.ld(self), ag__.ld(iterator)), None, fscope)
      16             except:
      17                 do_return = False

```

**ValueError:** in user code:

```

File "/usr/local/lib/python3.8/dist-packages/keras/engine/training.py", line
1845, in predict_function *
    return step_function(self, iterator)
File "/usr/local/lib/python3.8/dist-packages/keras/engine/training.py", line
1834, in step_function **
    outputs = model.distribute_strategy.run(run_step, args=(data,))
File "/usr/local/lib/python3.8/dist-packages/keras/engine/training.py", line
1823, in run_step **
    outputs = model.predict_step(data)
File "/usr/local/lib/python3.8/dist-packages/keras/engine/training.py", line
1791, in predict_step
    return self(x, training=False)
File "/usr/local/lib/python3.8/dist-packages/keras/utils/traceback_utils.py",
line 67, in error_handler
    raise e.with_traceback(filtered_tb) from None
File "/usr/local/lib/python3.8/dist-packages/keras/engine/input_spec.py", line

```

ψ

Next is the Embedding layer, the layer would contain a 5000 vocab for the embedding layer with 15 input length. After the layer would flatten the nodes and have a 3 layer 32, 8, 1 setup. Again, the embedding layer didn't do too well with only an accuracy of .51 in train and on the evaluation.

In all, due to the size of the dataset, the models couldn't shine as the DL require a large amount of data unfortunately. But the best model of them all is the sequential model.

Next is the Embedding layer, the layer would contain a 5000 vocab for the embedding layer with 15 input length. After the layer would flatten the nodes and have a 3 layer 32, 8, 1 setup. Again, the embedding layer didn't do too well with only an accuracy of .51 in train and on the evaluation.

In all, due to the size of the dataset, the models couldn't shine as the DL require a large amount of data unfortunately. But the best model of them all is the sequential model.



[Colab paid products](#) - [Cancel contracts here](#)